

Relevance of Medical Questions

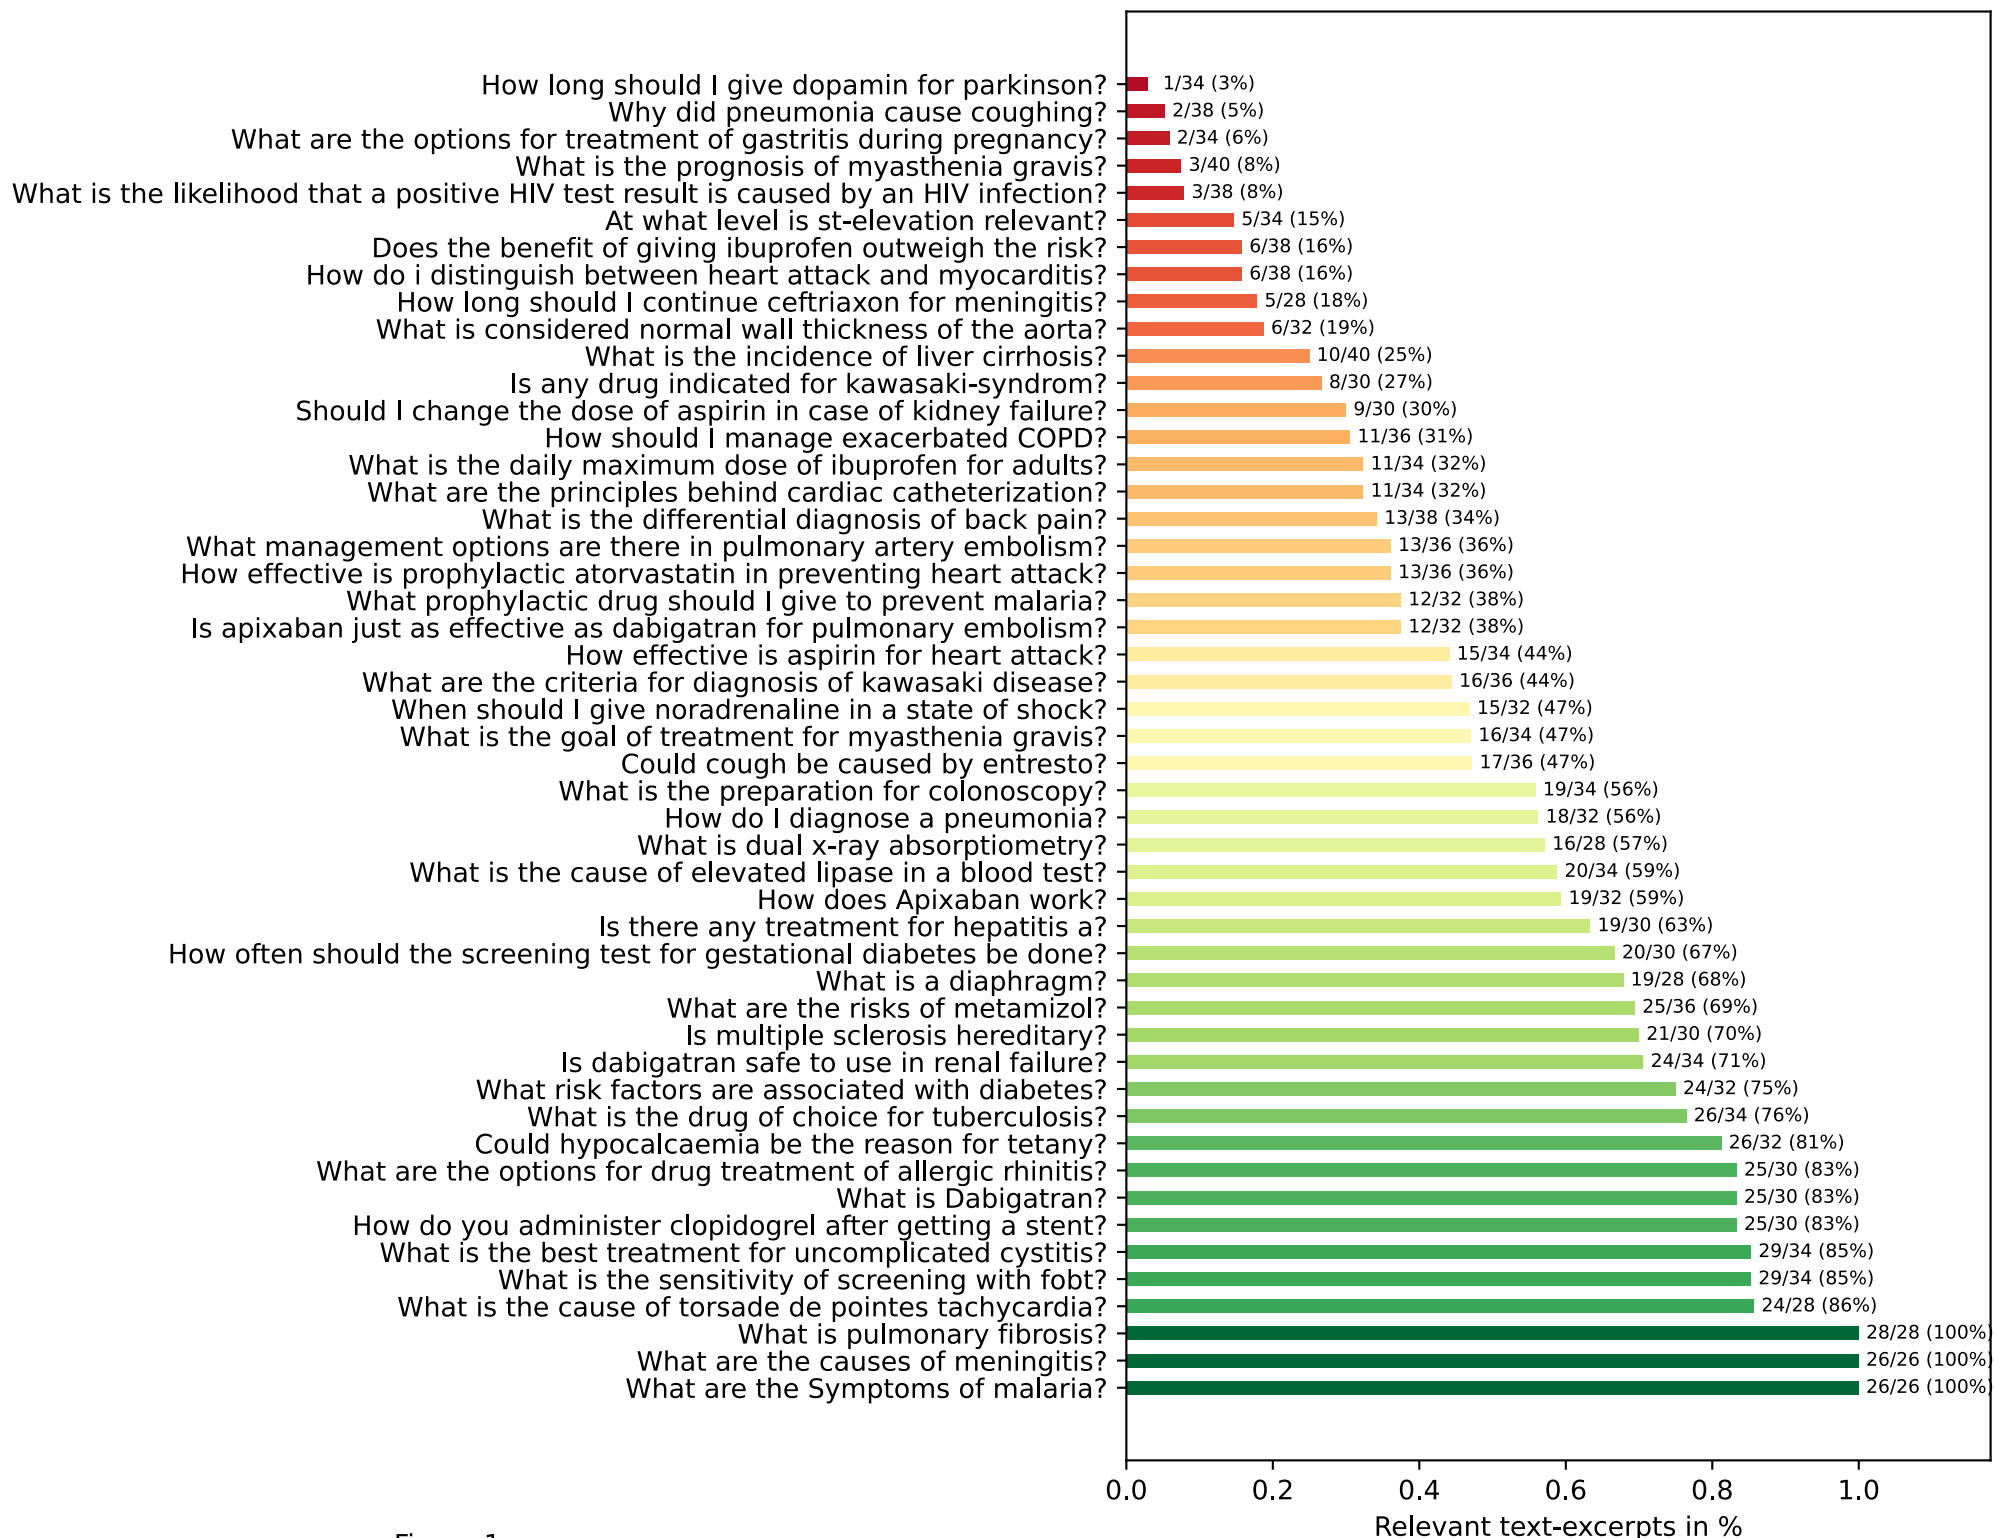


Figure 1:
 Relevance of Medical Questions.
 On Y-Axis are all queries.
 On X-Axis is the relevance in % per query.
 The figure shoos the distribution of relevance labels per query,
 showing the proportion of relevant labels
 annotations across the 50 clinical questions used in this evaluation.
 e.g. How long should I give dopamin for parkinson? There is 1 relevant - from 34 possible text-excerpts.
 Red means more irrelevant text-excerpts per query, and green means more relevant text-excerpts.

Batch A

Batch B

NDCG@10 per query — disjoint query sets (n=25 each)



Figure 2:
NDCG@10 across Batches.
Seperated Heatmaps per Batch
On Y-Axis are all queries.
On X-Axis is the model name.
Comparison of NDCG@10 scores between the two disjoint question batches
Batch A(left) and Batch B(right), showing the per-query ranking quality
for each of the four model variants.
Red means more irrelevant text-excerpts has been retrieved from the model per query,
and green means more relevant text-excerpts.

Label agreement

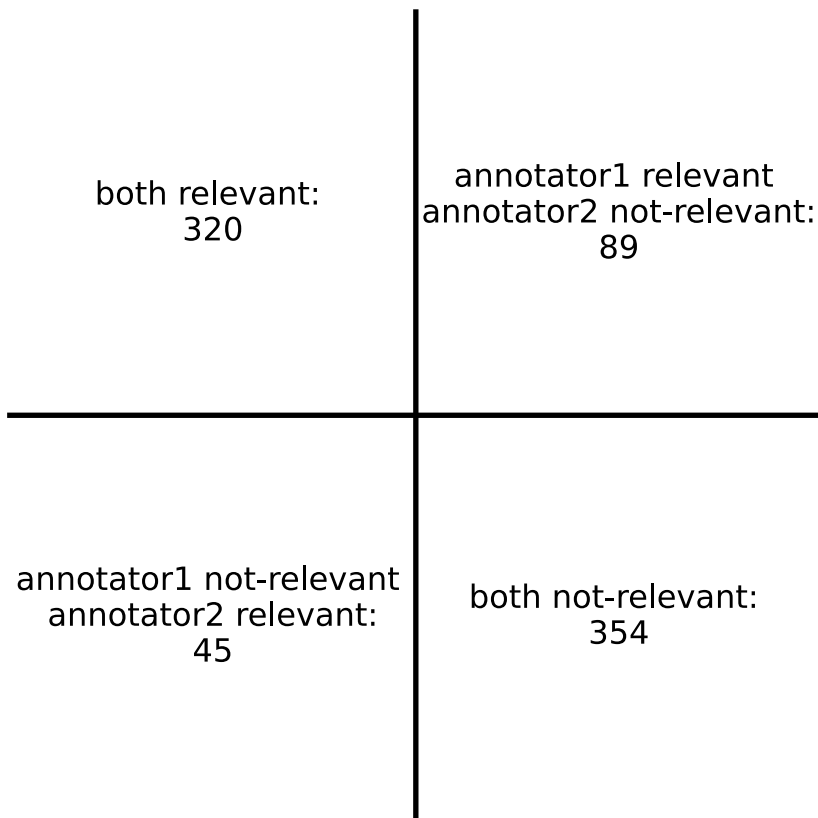


Figure 3:

Label agreement

The figure shows the agreement and disagreement of the annotators, for all annotated data.